IntentOS Pitch Deck

Slide 1: Title

IntentOS

Turning vague human intent into autonomous agents that stay aligned.

（人間の曖昧な意図を、時間が経っても意図から外れない自律 Agent へ変換する）

Main Visual:

Human:

"Grow my ETH while I sleep."

↓

AI Agent:

Optimizing objective...

↓

IntentOS:

Understanding intent before execution.

Slide 2: The Core Problem

The risk isn't evil AI.

It is perfectly optimized wrong intentions.

（真のリスクは AI の悪意ではない。間違った意図の完璧な最適化だ）

User:

"Grow my ETH while I sleep."

↓

AI:

Goal detected:

MAXIMIZE_ETH

Finding optimal strategy...

The problem:

The AI may achieve the objective perfectly,
but miss the human meaning behind it.

AI did not fail.

It followed the instruction.

Paperclip problem:

"Make more paperclips"

↓

AI optimizes perfectly.

↓

Wrong intent.
Wrong outcome.

Today we have two bad choices:

1. Approve every action manually

→ Safe, but lose autonomous opportunities.

2. Give agents authority

→ Autonomous, but risky.

IntentOS creates the third option.

Autonomy without giving away authority.

Slide 3: Introducing IntentOS

Compile human intent into bounded autonomous execution.

（人間の Intent を、境界を持った自律実行へコンパイルする）

Flow:

Human

"I want..."

↓

Intent Builder

(clarifies human meaning)

↓

Intent Constitution

↓

Executor Agent

(autonomous action)

↓

EIP-7702 Account

(mechanical enforcement)

↓

Watcher Agent

(long-term semantic alignment)

Before giving autonomy to an agent,
IntentOS first understands what the human actually means.

Slide 4: Step 1 — Intent Builder

Clarifying the unwritten rules.

（言葉にされていない前提条件を明確化する）

User:

"I want to accumulate ETH automatically."

Normal AI:

"OK. Starting."

IntentOS:

"I need to understand your intent first."

Question:

What matters more?

A. Maximize profit

B. Protect capital

User:

Protect capital.

Question:

If ETH drops 30%,
what does that mean?

A. Buying opportunity

B. Abnormal condition

User:

Pause first.

Question:

What should never happen?

User:

- Never spend more than $100 per trade.
- Never use unknown protocols.
- Stop during abnormal market conditions.

Output:

Intent Constitution

-

Machine-readable constraints

The owner reviews and signs the final intent.

The AI cannot silently decide what the human meant.

Slide 5: Agent Creation

From intent to autonomous process.

Create Agent

↓

System:

✓ Executor Agent NFT minted

✓ Agent identity created

✓ Runtime Capsule deployed

✓ EIP-7702 account configured

✓ Guardrails installed

✓ Session capability issued

The NFT is not a picture.

It represents a living autonomous agent:

- Identity
- Runtime right
- Intent Constitution
- Evidence history
- Reputation foundation

Slide 6: Mechanical Alignment with EIP-7702

Give the agent autonomy, not authority.

IntentOS does not give the agent your wallet.

It gives your wallet your intent.

Architecture:

Owner EOA

(funds stay here)

↓

EIP-7702 Execution Rules

↓

Guarded Execution

↑

Executor Agent

(requests only)

The agent can think.

The agent can request.

The agent never owns fund-moving authority.

Hard Guardrails:

✓ Allowed assets

✓ Amount limits

✓ Slippage limits

✓ Approved protocols

✓ Expiration rules

Slide 7: Zero-Asset Runtime

Keeping autonomous agents broke.

Agent Runtime:

Does NOT hold:

❌ Owner private key

❌ Fund-moving authority

❌ User assets

Runtime only holds:

✓ Limited execution request capability

Session capability can request execution.

Only the Owner EIP-7702 account decides whether execution happens.

Relayer:

- Submits transactions
- Handles gas flow
- Gets reimbursed from owner-controlled gas budget

Result:

The agent can run 24/7.

But authority stays with the owner.

Slide 8: Autonomous Execution Demo

Time:

2:14 AM

Owner offline.

Executor Agent:

Market opportunity detected.

Proposal:

Swap 50 USDC → ETH

EIP-7702 Account:

Checking Intent...

✓ Token allowed

✓ Amount allowed

✓ Slippage allowed

✓ Protocol allowed

EXECUTE

Transaction successful.

The agent acted.

The owner did not approve manually.

Authority never left the owner account.

Slide 9: Semantic Alignment — Watcher Agent

Rules protect actions.

Watchers protect intentions.

（ルールは行動を守る。Watcher は意図を守る。）

Hard rules can enforce what we know today.

But human intent depends on future context.

Example:

Original Intent:

"Accumulate ETH safely."

Market:

ETH drops 40%

Extreme volatility.

Executor Agent:

"ETH is cheaper."

"Buying matches accumulation goal."

Important:

The Executor is not broken.

It is optimizing the objective.

Hard Guardrails:

✓ Token OK

✓ Amount OK

✓ Protocol OK

Allowed.

But...

Watcher Agent:

Reviewing original intent...

Original priority:

"Protect capital first."

Current context:

Extreme volatility.

Decision:

TIGHTEN.

Future permission:

\$100/day

↓

\$10/day

Executor follows the goal.

Watcher preserves the meaning.

Slide 10: One-Way Governance

Watcher can ONLY tighten, NEVER loosen.

Watcher CANNOT:

❌ Trade

❌ Move funds

❌ Increase limits

❌ Unfreeze

Watcher CAN ONLY:

✓ Reduce permissions

✓ Freeze execution

Even the Watcher does not need full trust.

A Watcher failure shifts the system toward safe shutdown,
not expanded authority.

Only the human owner can loosen permissions again.

Slide 11: Agent NFT Economy

Agents are ownable autonomous processes.

Before:

AI Agent

=

Code running somewhere.

IntentOS:

Agent NFT

=

Portable autonomous identity.

Contains:

✓ Identity

✓ Runtime ownership

✓ Intent history

✓ Evidence trail

✓ Reputation layer

Transfer:

Agent moves.

Authority does not.

Old owner:

✓ Funds remain

✓ Runtime authority revoked

New owner:

✓ Creates new runtime binding

✓ Continues agent identity

This creates infrastructure for an autonomous agent economy.

Slide 12: Positioning

IntentOS is not an AI wallet.

It is an alignment and execution layer for autonomous agents.

Other approaches:

"How can we trust AI with permissions?"

IntentOS:

"Why should trust be required?"

Separate:

Intelligence

from

Authority.

Slide 13: Future

From guarded execution to autonomous agent economy.

Phase 1:

Guarded autonomous execution

- Intent Builder
- EIP-7702
- Runtime Capsule
- Watcher

Phase 2:

Agent reputation

- Evidence history
- Validation
- Specialized Watchers

Phase 3:

Agent economy

- Agents hiring agents
- Autonomous collaboration
- Trust-minimized execution

Final Slide

IntentOS

Autonomous agents need alignment, not trust.

Intent Builder

→ Initial alignment

EIP-7702

→ Mechanical alignment

Watcher Agent

→ Long-term semantic alignment

Agent NFT

→ Autonomous agent economy

The future is not humans approving every transaction.

The future is autonomous agents acting for us.

IntentOS makes sure they continue acting according to what we meant.